

Julian Kaljuvee
kaljuvee@gmail.com – +44 7585 099 329 (m)
<https://kaljuvee.github.io>

SUMMARY

- Senior technology, data and risk leader with over ten years of experience in statistics, data science, and software / data / ML engineering, primarily in financial services, private equity (PE), semiconductor, retail / FMCG domains
- Strong knowledge of statistical methods, machine learning, artificial intelligence, and general software engineering techniques across a wide set of technologies and programming languages
- Professional work / consulting experience includes UBS, London Stock Exchange, Standard Chartered, HSBC, DBRS Morningstar, Nandos UK, Vaxart, Microsoft and Indurent

DATA and SOFTWARE ENGINEERING

Languages:	Python, R, Java/Scala/Groovy
Cloud:	GCP, Azure, AWS
Databases/Big Data:	PostgreSQL, Presto, Athena, Redshift, BigQuery, MongoDB, Databricks, Snowflake, PySpark/SQLAlchemy, MS-SQL, Neo4J
DevOps:	Kubernetes, Docker, Terraform, Pulumi
ML/AI libraries:	PyTorch, JAX, PyCaret, SciKit-Learn, spaCy, FB Prophet, Transformers; Foundational LLMs: OpenAI, Anthropic, LLaMA, Mistral; Orchestration: LlamaIndex, Semantic Kernel, PromptFlow, AutoGen, LangGraph; Reinforcement Learning (RL): Q-learning (Time temporal, Monte Carlo), Deep Q-Network (DQN), Proximal Policy Optimization (PPO)
Visualisation:	Plot.ly, Streamlit, Tableau, PowewBI, Looker, Metabase

PROFESSIONAL EXPERIENCE

Indurent (www.indurent.com) (Blackstone) London, UK / Remote Sep 2025 - present

AI Engineer (Contract)

Business: AI engineer to lead several key AI initiatives at Indurent, a £6bn+ Blackstone portfolio company in PE / CRE space. Responsibilities include supporting platform scale-up from a £700m REIT, supporting a multi-billion-pound merger integration, and embedding production AI capabilities across core business and investment functions

Responsibilities: (i) AI Lease abstraction - processing tenant leases automatically, (ii) AI Invoice abstraction - automatically processing and labelling invoices for EPC categories, (iii) AI risk scoring - developing and implementing automated scoring methodology for technical and environmental risk from based on due diligence reports, (iv) Comparable analysis and automated valuation models (AVM) - analysing and valuing property assets in an automated manner

Stack: Azure AI Services, Fabric / DBT, Streamlit / FastHTML UI; Azure AI Search and AI Foundry, Agentic AI (LangGraph/Langchain), model evaluation (DeepEval/DSpy); Knowledge Graphs, GraphRAG, Neo4J, Azure CosmosDB (Gremlin API)

Microsoft (NASDAQ:MSFT) (www.microsoft.com) London, UK / Remote Apr 2024 - Sep 2025

Gen AI Engineer (Contract)

Responsibilities: generative AI solutions for a number of enterprise clients for large Fortune 100 global clients across different industries (mining, automotive, oil / gas). Use cases have ranged from RAG optimisation, AI search, financial advisory, customer journey predictive analytics (call transcription analysis for energy retail sales), image classification / computer visualization to AI reasoning agents with real-time data via enterprise APIs. Clients included [ADNOC \(Abu Dhabi National Oil Company\)](#), [Mubadala](#), [Anglo American](#), [Ternium](#), [Bosch](#) and [EDF Energy](#)

Observability / evaluations: Implemented AI observability and safety tools such as Guardrails AI, LangFuse, and LangSmith

Stakeholder management / leadership: engaged project stakeholders for requirements analysis, and AI strategy definition to ensure customer success. As a team lead, mentored and guided a team of data scientists / AI engineers, including upskilling and knowledge sharing.

Stack: Azure AI Services, Cosmos DB, Snowflake, Databricks; Streamlit UI; Azure AI Search and AI Foundry, GenAI, LLM finetuning, Agentic AI (LangGraph/Langchain, MSFT Agent Framework, Magentic-One), model evaluation; Vector DBs: Pinecone, Weaviate, ChromaDB; Protocols: MCP, A2A; MLOps/LLMOps: MLFlow

Vaxart, Inc (NASDAQ:VXRT) (www.vaxart.com) San Francisco, US / Remote Jun 2023 - Mar 2024

Head of AI (Contract)

Business Area: biotech, AI-driven drug design, manufacturing / bioreactor process optimisation and clinical trial optimisation

Responsibilities: (i) defined and implemented a proprietary platform to optimise vaccine bioreactor manufacturing process, including custom LLMs for analysing data (RAG), (ii) build a POC for Data ingestion, In-silico mutagenesis experiments, Cross validation and molecular dynamics and Mutation probabilities sandbox leveraging DeepChain.bio platform, (iii) LLM-based approaches for optimising clinical trial design, monitoring and regulatory reporting.

Stack: GCP / BigQuery, Snowflake, Databricks, Streamlit; Python, GenAI: LangChain, OpenAI, Anthropic, LlamaIndex; Snowflake

Worldremit, (www.worldremit.com) London, UK / Remote Oct 2022 - Jun 2023

Head of Data Engineering (Contract)

Business Area: revenue / sales forecasting, cloud migration, reconciliation of EOD and EOM finance activities for a global Fintech

Responsibilities: (i) productionising monthly revenue forecasting / predictive analytics engine (MLFlow, SageMaker), (ii) development and deployment of new global finance reconciliation engine against 200+ FX correspondents globally with over 3m monthly transactions in a multi-cloud environment (DBT/Redshift), (iii) calculating KPIs for CFO office such as monthly active users, and cohort analysis

Data modelling and architecture: defined target data model and architecture for the new greenfield data warehouse that hosted millions of records monthly payment data

Stack: Git / Github, Airflow; Metabase; AWS, GCP / BigQuery, DBT and Redshift; Databricks/PySpark, Prophet, MLFlow, Python

IKEA, (www.ikea.com)

Malmö, Sweden / Remote

Jan 2022 - Oct 2022

Senior Data Scientist / Engineer (Contract)

Business Area: group wide KPIs and business strategy planning for IKEA / INGKA group

Responsibilities: supporting development and deployment of new data lake / data warehouse (DWH) for reporting group wide metrics. Data included financial, sustainability, people, customer-related metrics (including Google Analytics and AdWords APIs), online customer journey optimisation. Dashboards in Tableau and Looker for stakeholders to display KPI data

Data architecture and governance: defined data model, standards and governance for the KPI data warehouse (GDPR out of scope as data were aggregated and masked upstream)

Stack: Git / Github, Airflow; Tableau / Looker / PowerBI; Python (Pandas, Scikit-learn), Apache Airflow / DBT; GCP / BigQuery / DataFlow / DataForm; Terraform / Pulumi (IaC)

Nandos, (www.nandos.co.uk)

London, UK

Jun 2020 - Dec 2021

Senior Data Scientist / ML Engineer (Contract)

Business Area: sales / demand forecasting, data infrastructure development / data lake, POS, loyalty, rota, weather, wifi and geolocation and delivery data integration for a major restaurant chain

Responsibilities: supporting development and deployment new data lake / data warehouse (DWH) infrastructure (billions of historical records), sales forecasting (time series analysis), marketing (Google Ads and AdWords APIs with Budgeting and Campaigns), and customer journey prediction and loyalty data visualization tools; implementing ETLs on batch data, CI/CD best practices (Airflow and Github actions), documentation, code refactoring, mentoring and training. Dashboards in Tableau and Looker for stakeholders to measure KPI data

Marketing mix model (MMM) development: developed media marketing mix model (MMM) using the classical regression model (base sales, coefficients of media spend and pricing / non pricing promotions). Validated and optimised model against core KPIs (CAC, market share, sales revenue, ROAS) and supported productionisation of the model by the marketing team

Stack: Git / Github, Airflow; Tableau / Looker / Superset; Python (Pandas, Scikit-learn, FB Prophet), Apache Airflow / DBT; GCP / BigQuery, AWS (EC2, S3, Athena, Redshift, Presto), deployment Kubernetes (EKS), Databricks

DBRS Morningstar, (<http://www.morningstar.com>)

London, UK / Frankfurt, DE

Jul 2019 - Jun 2020

Quantitative Researcher / Data Scientist (Contract)

Business Area: structured finance and credit ratings

Responsibilities: supporting development and deployment of EU RMBS rating models for a leading global financial analytics and ratings provider integrating with European Data Warehouse (over 2bn records)

Model building and deployment: (i) data analysis / visualisation, (ii) data cleansing and transformation (data sets with over 100m records), (iii) model development, feature engineering and model evaluation with statsmodels, Logistic Regression, GAM and XGBoost, (iv) model deployment and productionisation with MLFlow (v) presentations to business stakeholders

Stack: Git / Github, Domino Datalab, Jupyter, Jenkins; Python (Pandas, Scikit-learn, Statsmodels, Plotly, PySpark), MLFlow, Docker, Kubernetes, AWS (S3, Athena, Redshift), Databricks

Quantitative Analyst / Data Scientist- Financial Services (Contractor)

UBS, Data Scientist, Market risk management (Jun 2018 - Jul 2019)

HSBC, Data Scientist / Quantitative Analyst, Market risk management (Jan 2018 - Jun 2018)

Standard Chartered Bank, Data Scientist / Quantitative Analyst, Market risk management (Mar 2016 - Dec 2017)

HSBC, Quantitative Analyst, Product Control (Nov 2015 - Feb 2016)

London Stock Exchange Group (LSEG) / LCH, Data Scientist, Market risk (Jul 2014 - Nov 2015)

UBS, Quantitative Analyst, IRD/FX/MM Front-Office (Jan 2010 - Jun 2014)

EDUCATION

Harvard University, Cambridge, Massachusetts
BS in Applied Mathematics and Economics

Course work included Economics, Applied Mathematics, Computer Science, Molecular Biology and Physics

Columbia University, New York, New York

MS level studies in Statistics and Probability Theory

Studies included Statistics, Probability Theory and Operational Research

TALKS

- [Data Scrum Alternative Data - Sep 2019](#)
- [Estonia AI Meetup - Jul 2023](#)

- [Estonia AI Meetup - Oct 2023](#)
- [Mentastic AI Workshop - Sep 2024](#)
- [Estonia AI Meetup - Mar 2025](#)
- [RAG Enterprise Search - Jun 2025](#)

PRO

Startups: co-founder of and advisor to several VC backed Fintech startups, including [Bondora.com](#), [EstateGuru.co](#) and [Investly.co](#)

Projects: [PE Hero](#) - financial research / simulation platform

Github: <https://github.com/kaljuvee/> and <https://kaljuvee.github.io/>

TECH SKILLS MATRIX

Languages

Python (Pandas, Numpy, SciPy)	5+ Years
Java/ Groovy / Scala	5+ years

ML / AI / Data Science

Machine Learning / Deep Learning: PyTorch/ JAX scikit-learn, Prophet	5+ Years
GenAI: LLMs - OpenAI, Anthropic, Gemini, LLama; LangChain, HuggingFace; Orchestration: LangGraph, Autogen, Crew AI, Semantic Kernel Evaluation: Ragas, DeepEval, TrueLens, DSPy Protocols: MCP, A2A	3 Years 1 year
AutoML: AWS SageMaker, Azure ML, Google Vertex.AI	5 Years

Visualisation

Matplotlib, Seaborn, Plot.ly, Tableau, Apache Superset, Power BI, Looker	5 Years
---	---------

Devops / Data Scientisting / Dev Tools

Airflow	5 Years
Terraform	5 Years
Jenkins	5 Years
Docker, Kubernetes	5 Years

Cloud Technologies

Google Cloud Platform / GCP	3 Years
MS Azure	4 Years
AWS (EMR, Athena, S3)	4 Years

Middleware / Messaging

Kafka, Kinesis, PySpark / Databricks	4 Years
--------------------------------------	---------

Databases / ETL

Snowflake / Databricks / Redshift / BigQuery	5 years
RDMBS: MySQL, PostreSQL, Presto, MS SQL Server	> 10 Years